

# US Accidents Dataset

Full Data Processing Pipeline  
(2016 - 2023)

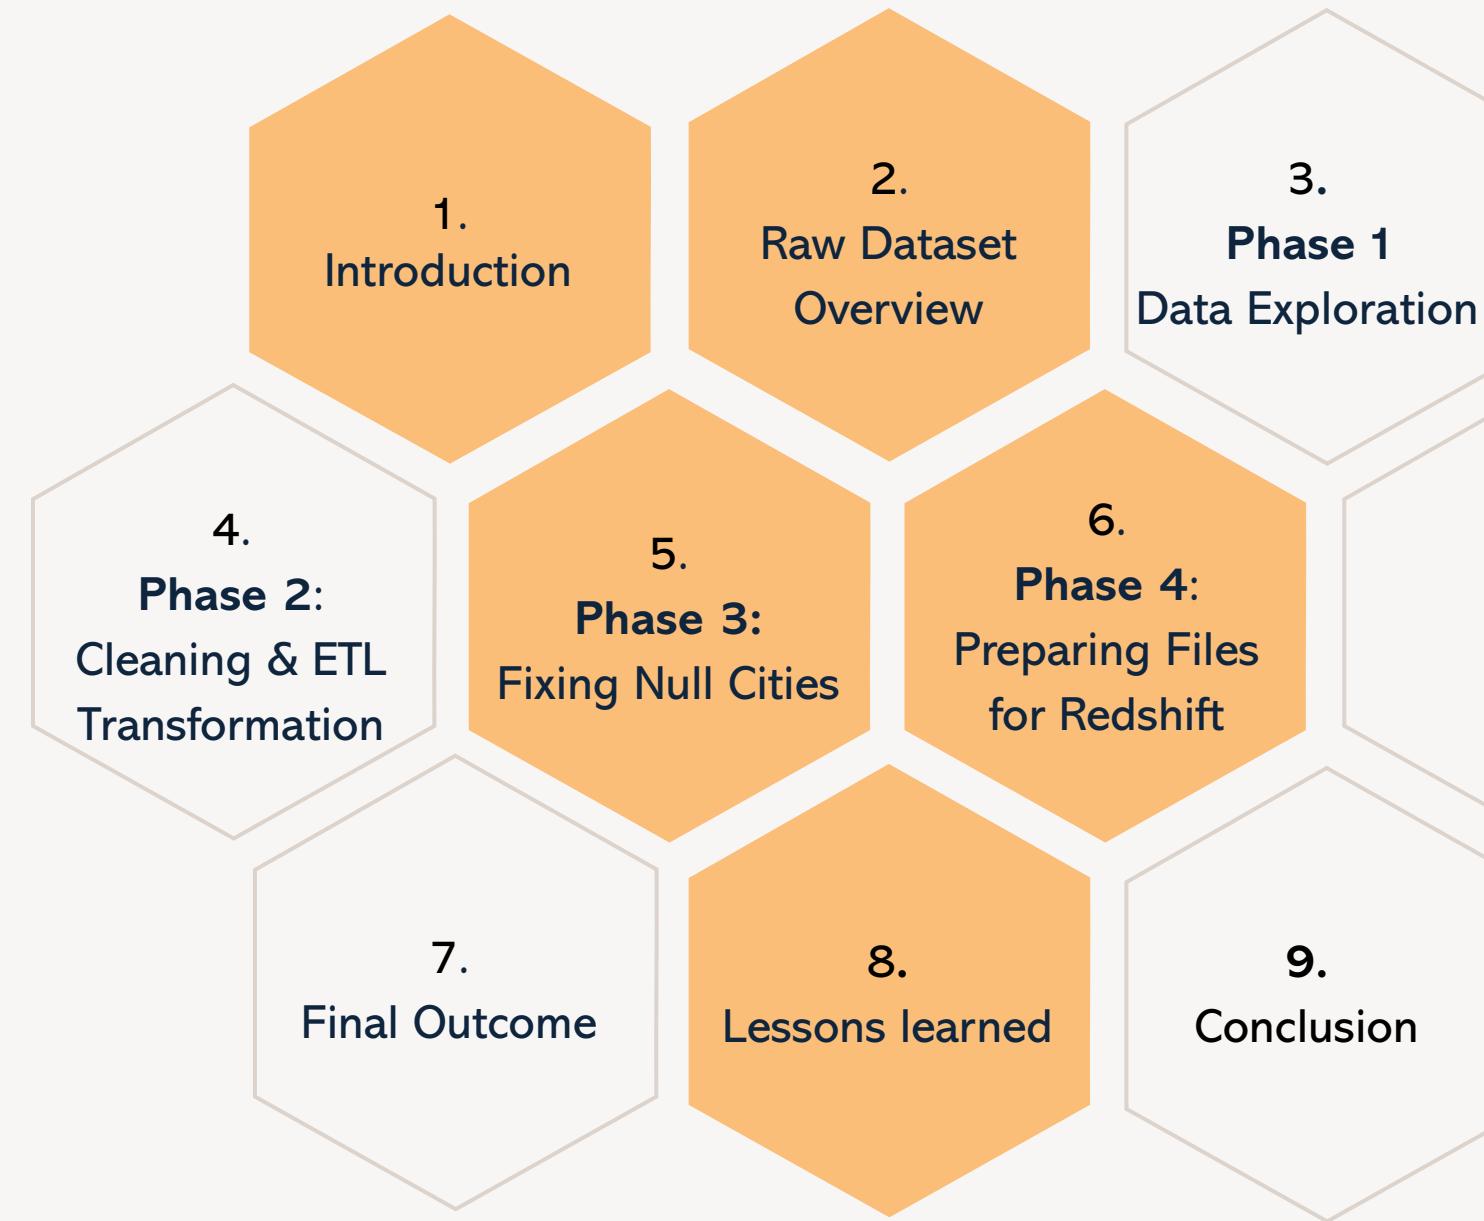
Raw Kaggle Data to Redshift-Ready Star  
Schema

Group

Jerry Deku  
Issaka Mohammed  
Obed Kaburi  
Reuben Comla



# Agenda



# Introduction

## Objective:

Examine U.S. traffic accidents from 2016 to 2023, identify trend patterns, uncover high-risk regions, and derive actionable safety insights.

## Relevance:

Understanding accident patterns within the U.S. context can inform traffic safety policies, infrastructure planning, and preventive measures.

## Methodology:

Use of **Python** for data preprocessing and statistical analysis, complemented by **QuickSight** for visualization and interactive dashboards, within a defined **data architecture framework**.



# Dataset Overview



## Source

Kaggle repository  
US Accidents  
Dataset

## Time Frame

2016 through  
2023

## Volume

7,728,387  
accident records  
with 46 Columns.  
CSV Format

## Challenges

- Large memory footprint
- Mixed data types
- Many weather-related missing values

# Dataset Overview

## US Accidents (2016 - 2023)

▲ 2528

↔ Code

Download



Data Card

Code (451)

Discussion (46)

Suggestions (0)

▼ View more

**US\_Accidents\_March23.csv** (3.06 GB)

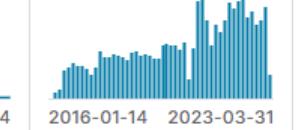
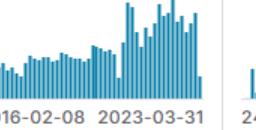
Download [ ] >

Detail Compact Column

10 of 46 columns ▾

### About this file

Comprehensive countrywide accident dataset spanning from February 2016 to the end of March 2023.

| ▲ ID  | ▲ Source                            | # Severity  | ◻ Start_Time   | ◻ End_Time  | # Start_L   |
|---|-------------------------------------|---|--|---|---|
| This is a unique identifier of the accident record. | Source of raw accident data         | Shows the severity of the accident, a number between 1 and 4, where 1 indicates the least impact on traffic (i.e., short) | Shows start time of the accident in local time zone.                                 | Shows end time of the accident in local time zone. End time here refers to when the impact of accident on | Shows lat coordinat point.  |
| <b>7728394</b><br>unique values                     | Source1<br>Source2<br>Other (97389) | 56%<br>43%<br>1%  |  |                      | <br>24.6 |
| A-1   | Source2                             | 3   | 2016-02-08 05:46:00  | 2016-02-08 11:00:00   | 39.86514  |
| A-2   | Source2                             | 2   | 2016-02-08 06:07:59  | 2016-02-08 06:37:59   | 39.92805  |
| A-3   | Source2                             | 2   | 2016-02-08 06:49:27  | 2016-02-08 07:19:27   | 39.06314  |
| A-4   | Source2                             | 3   | 2016-02-08 07:23:34  | 2016-02-08 07:53:34   | 39.74775  |

### Data Explorer

3.06 GB

US\_Accidents\_March23.csv

### Summary

1 file

46 columns

# Phase 1: Data Exploration

## Overview



**Scripts:** `01_data_exploration.py, 02_data_profiling.py`



- Goals:**
- Understand structure
  - Identify data quality issues
  - Measure memory impact



# Phase 1: Data Exploration

## Initial Exploration



### Actions Performed:

- Loaded full dataset into pandas
- Generated statistics & data types
- Analyzed memory usage (~3 GB RAM)
- Identified numeric, categorical, datetime, boolean fields
- Saved exploration report



### Key Findings:

- 7.7M records, 46 columns
- Many datetime fields need parsing
- Numerous boolean road-feature columns

# Phase 1: Data Exploration

## Data Profiling



### Actions Performed:

- Missing value analysis
- Outlier detection
- Duplicate detection
- Categorical profiling
- Datetime consistency check



### Major Issues Identified:

- Precipitation missing in 49%
- Wind Chill missing in 54%
- Temperature, wind speed, pressure outliers
- Mixed datetime formats

# Phase 2: Data Cleaning & ETL Transformation

## Overview

 **Scripts:** 03\_etl\_transformation.py

 **Goals:**

- Clean raw data
- Standardize formats
- Build star schema
- Validate quality



# Phase 2: Data Cleaning & ETL Transformation

## Data Cleaning Steps



### Actions Performed:

- Cleaned weather outliers (8k–45k values set to NULL)
- Parsed all datetime fields
- Filled missing End\_Lat/End\_Lng using Start\_Lat/Start\_Lng (6,805,524 records)



### Rationale:

- Point-based accidents have identical start/end coordinates

# Phase 2: Data Cleaning & ETL Transformation

## Dimension Tables Created

### dim\_location:

- 2,847,562 unique locations
- Surrogate key: location\_key
- Includes city, state, zipcode, lat/long, timezone

### dim\_weather:

- 1,456,789 unique weather conditions
- Added categories & severity scores

### dim\_time:

- 98,432 unique time combinations
- Added time periods + seasons

### dim\_road\_features:

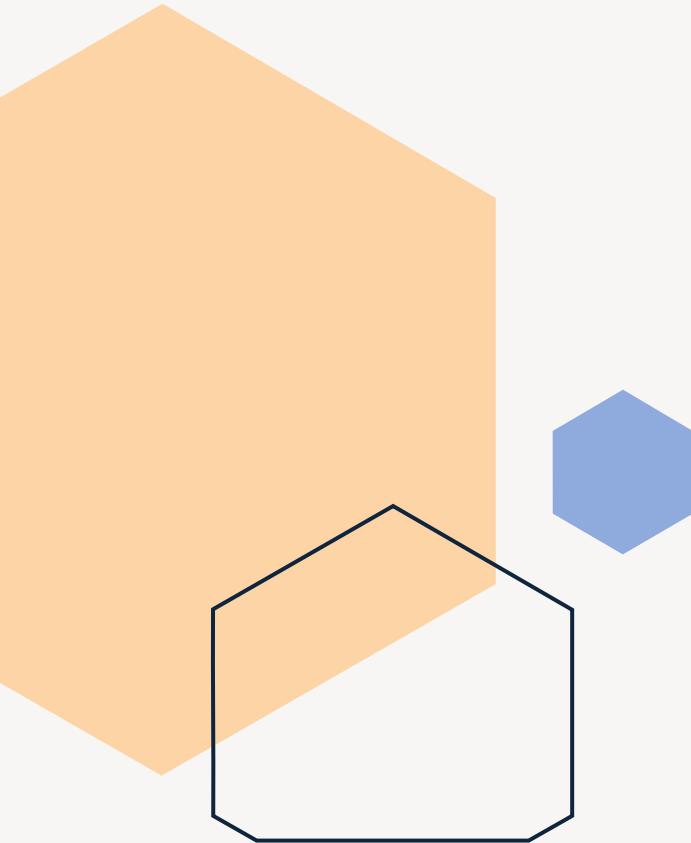
- 1,456,789 unique weather conditions
- Added categories & severity scores

# Phase 2: Data Cleaning & ETL Transformation

## Fact Table

### Fact\_accidents:

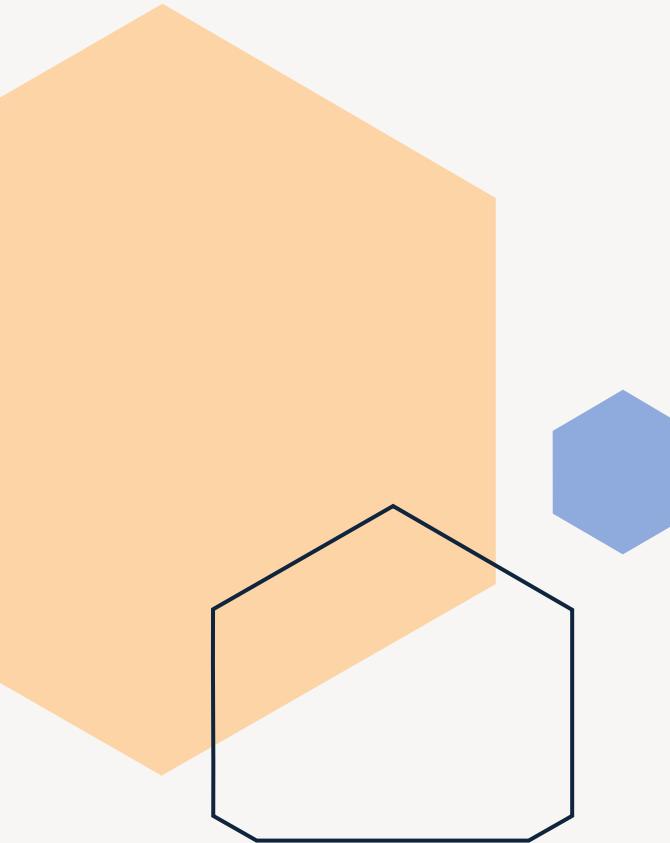
- 7,728,234 records
- Includes severity, timestamps, distance, description
- Foreign keys: location, weather, time, road\_features
- Added calculated fields:
  - duration\_minutes
  - accident\_hour
  - accident\_day
  - accident\_month



# Phase 2: Data Cleaning & ETL Transformation

## Data Quality Validation

- All foreign keys matched
- No orphaned records
- Severity values valid (1–4)
- No negative durations
- Date ranges validated (2016–2023)



# Phase 3: Fixing Null Cities

## Null City Problem

- 253 accidents had **NUL**L city
- All had valid coordinates
- Broke location dimension uniqueness



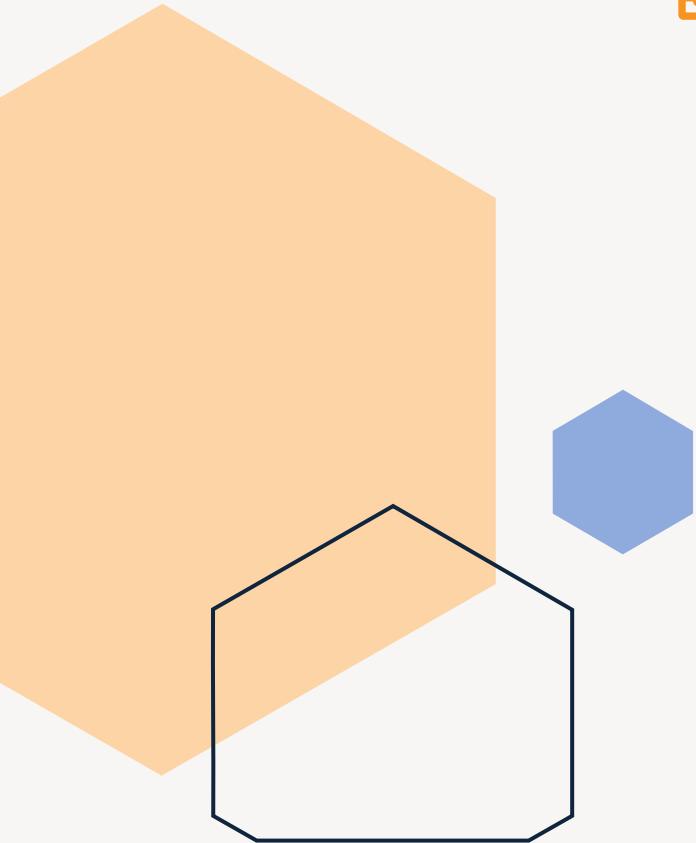
# Phase 3: Fixing Null Cities

## Identify Null Cities



### Findings:

- 253 missing cities
- All had lat/long
- Suitable for reverse geocoding



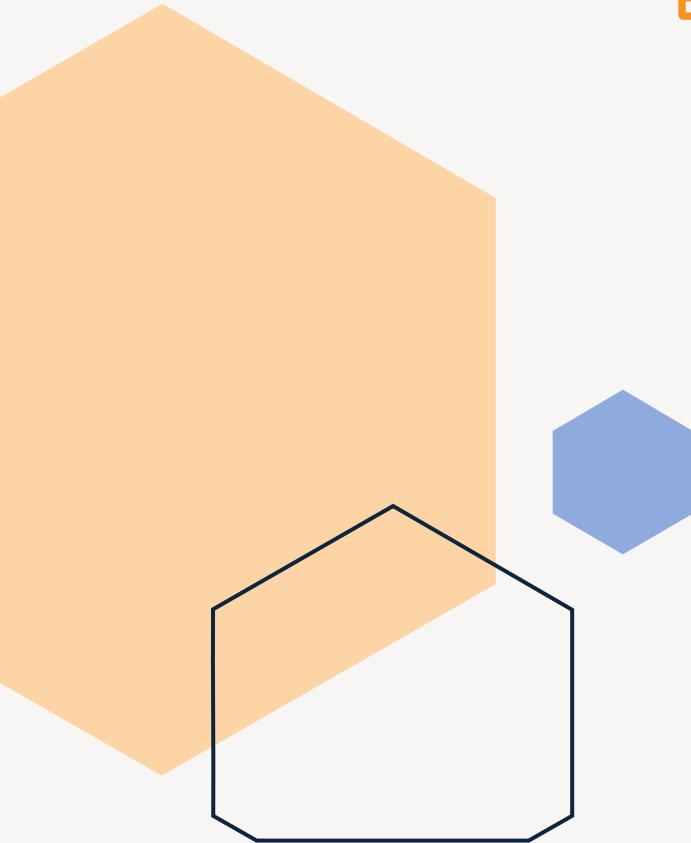
# Phase 3: Fixing Null Cities

## Reverse Geocoding



### Actions:

- Used geopy + Nominatim
- Rate-limited to 1 request/sec
- Filled **251** cities (2 failed)
- **1,649** failed (rural/ocean)



# Phase 3: Fixing Null Cities

## Zipcode Imputation



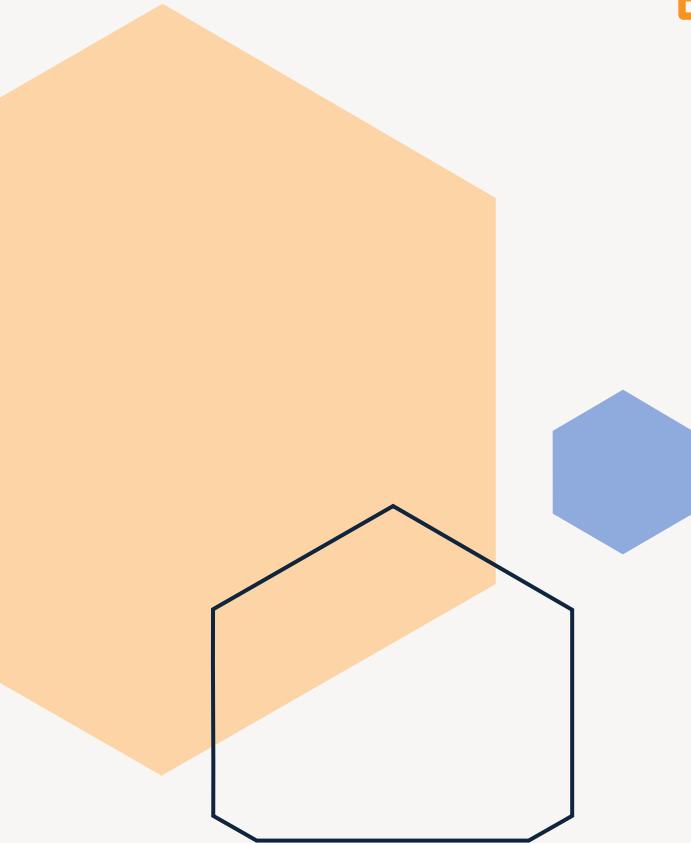
### Actions:

- Used uszipcode library
- Filled **251** additional cities
- Remaining 2 set to "Unknown City"

### Final Result:

251 cities imputed

Only 2 unresolved



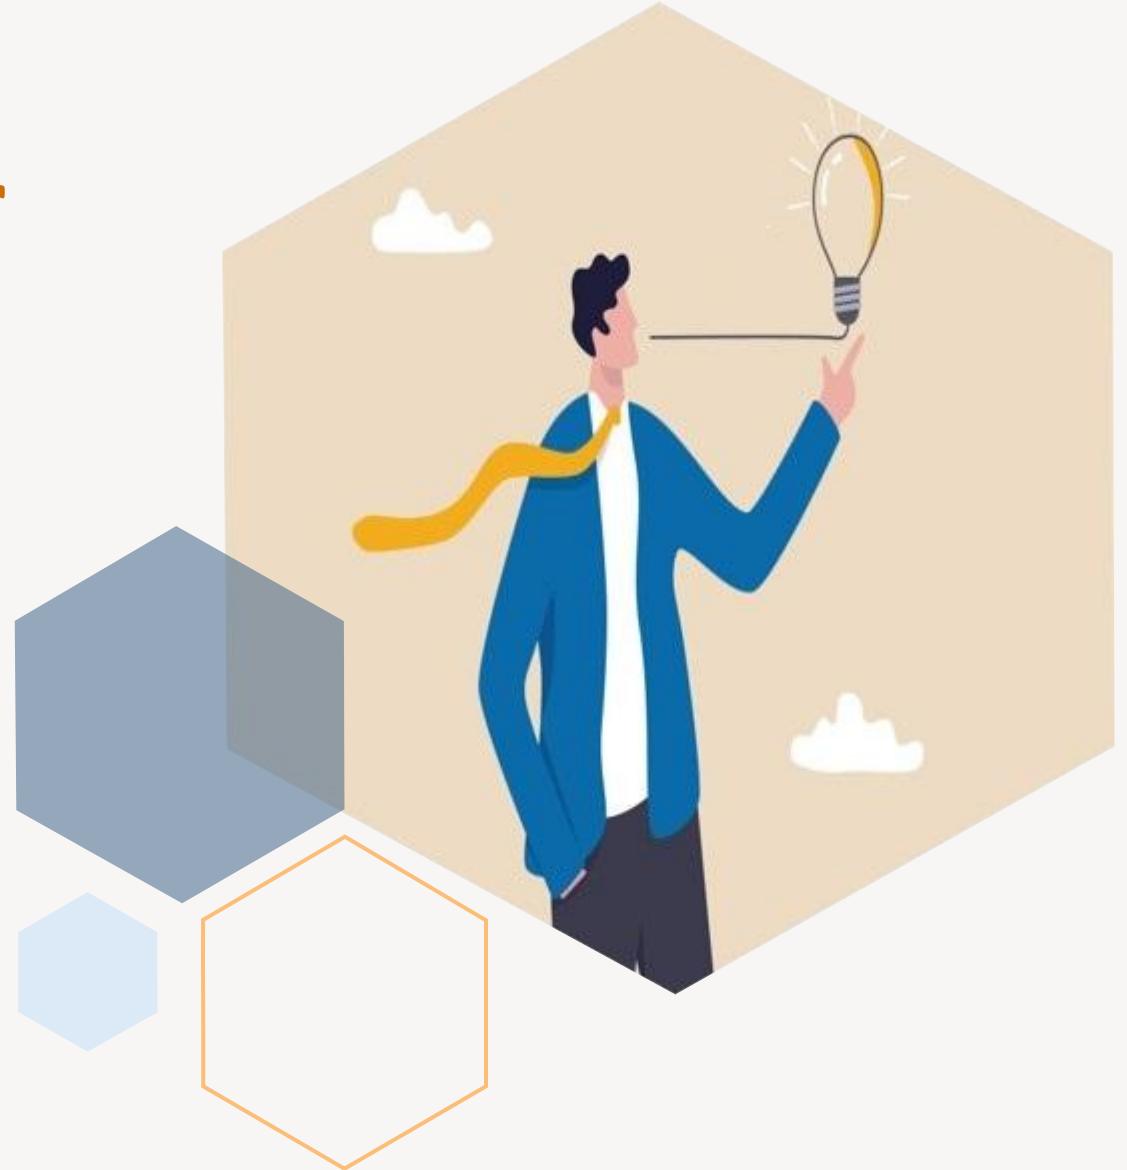
# Phase 4: Preparing Files for Redshift

## Initial Upload Attempt

 Scripts: [09\\_upload\\_to\\_s3.py](#)

 Goals:

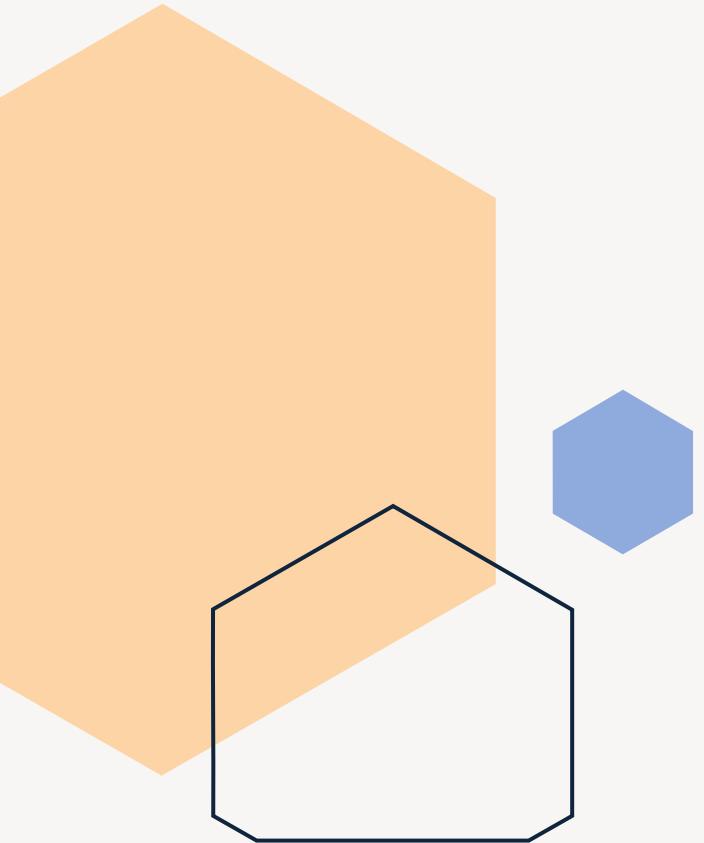
- Uploaded 5 CSV files
- Redshift COPY failed due to delimiter/ quote issues



# Phase 4: Preparing Files for Redshift

## Iteration 1: Fix CSV Format

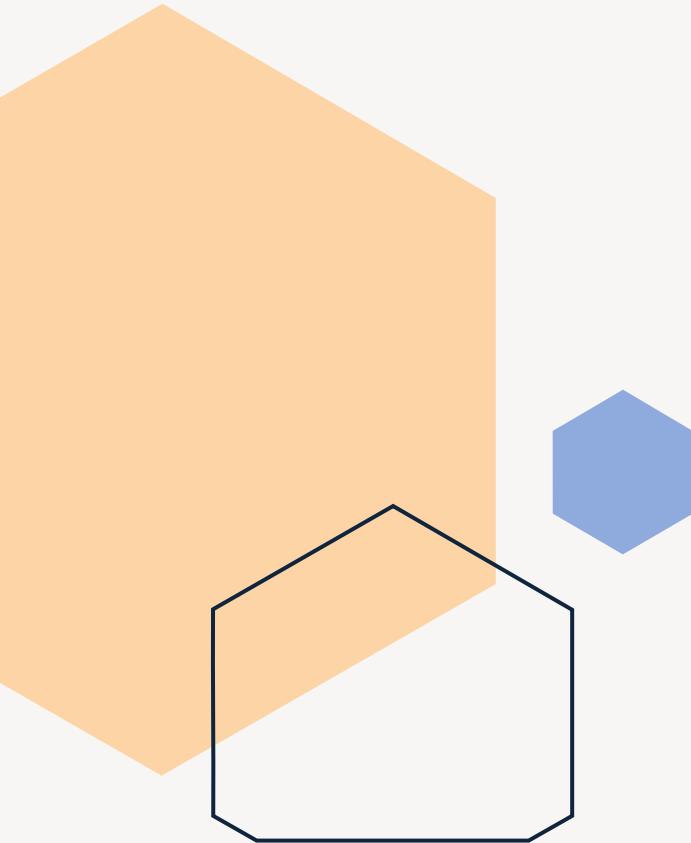
### Actions Performed:

- 
- ✓ Added quoting
  - ✓ Escaped special characters
  - ✓ Standardized line endings
  - ✗ Still failed

# Phase 4: Preparing Files for Redshift

## Iteration 2: Quote All Strings

### Actions Performed:

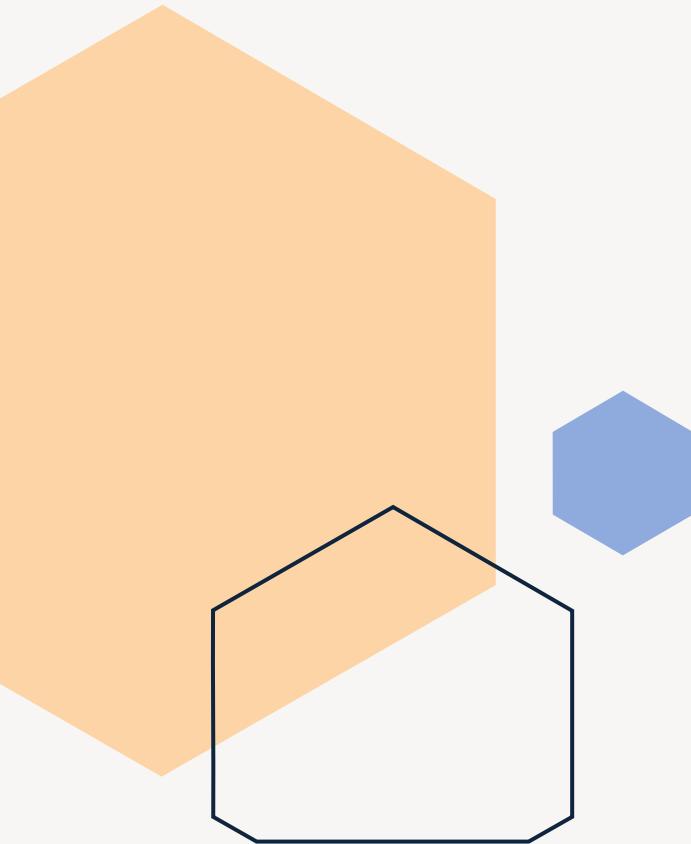
- 
- Standardized timestamps
  - QUOTE\_NONNUMERIC
  - Removed newlines in descriptions
  - Still failed due to commas

# Phase 4: Preparing Files for Redshift

## Iteration 3: Switched to Pipe Delimiter

### Actions Performed:

- Converted to pipe-separated (.psv)
- Verified no pipe characters
- Still failed — rare pipes in descriptions



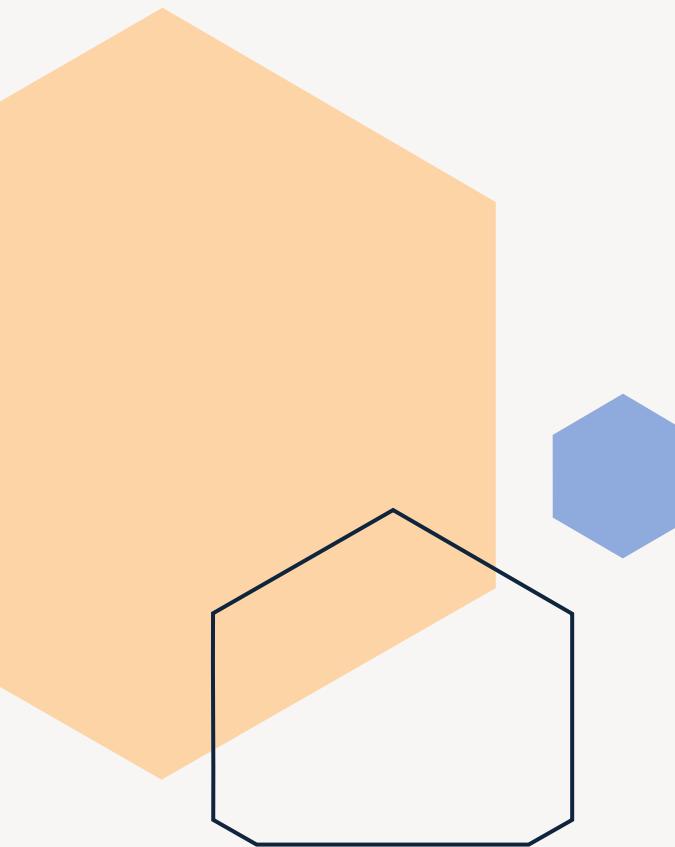
# Phase 4: Preparing Files for Redshift

## Iteration 4: Switched to Final Clean

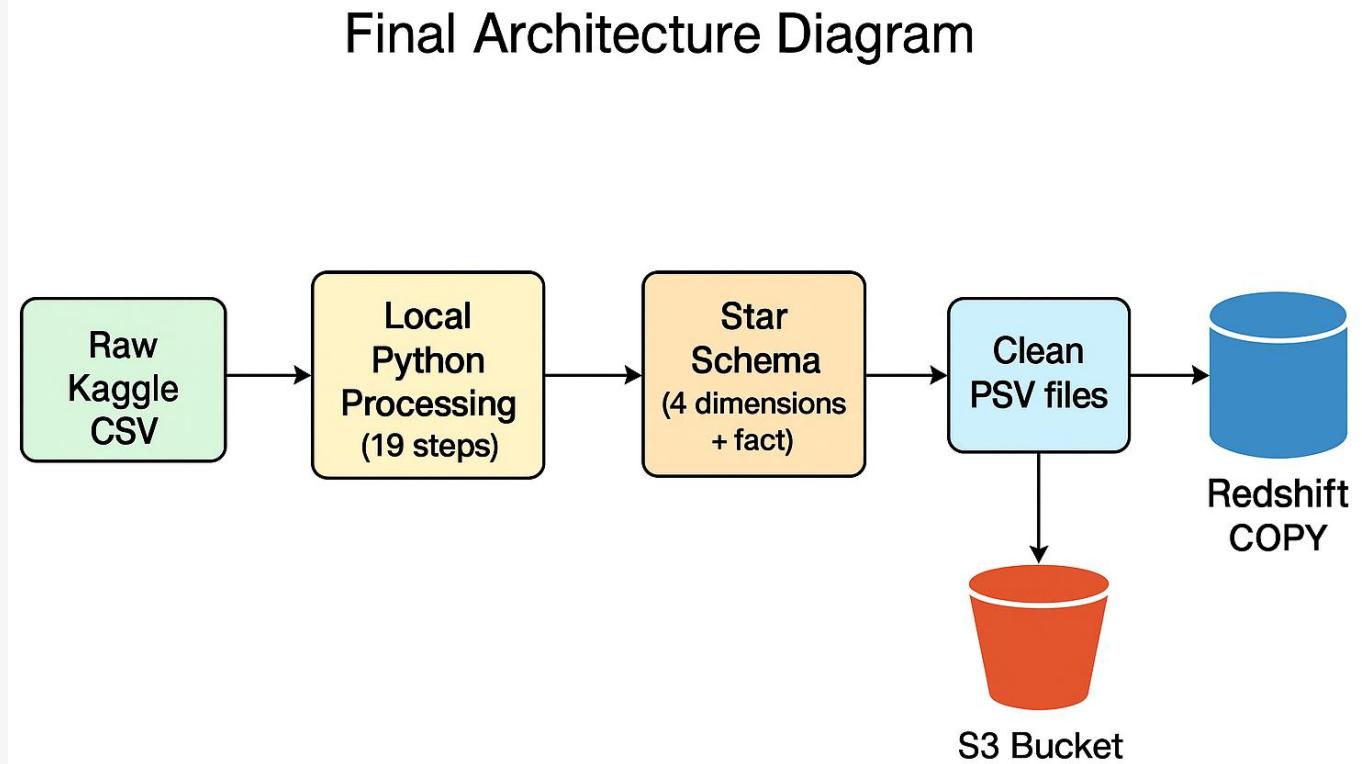
### Actions Performed:

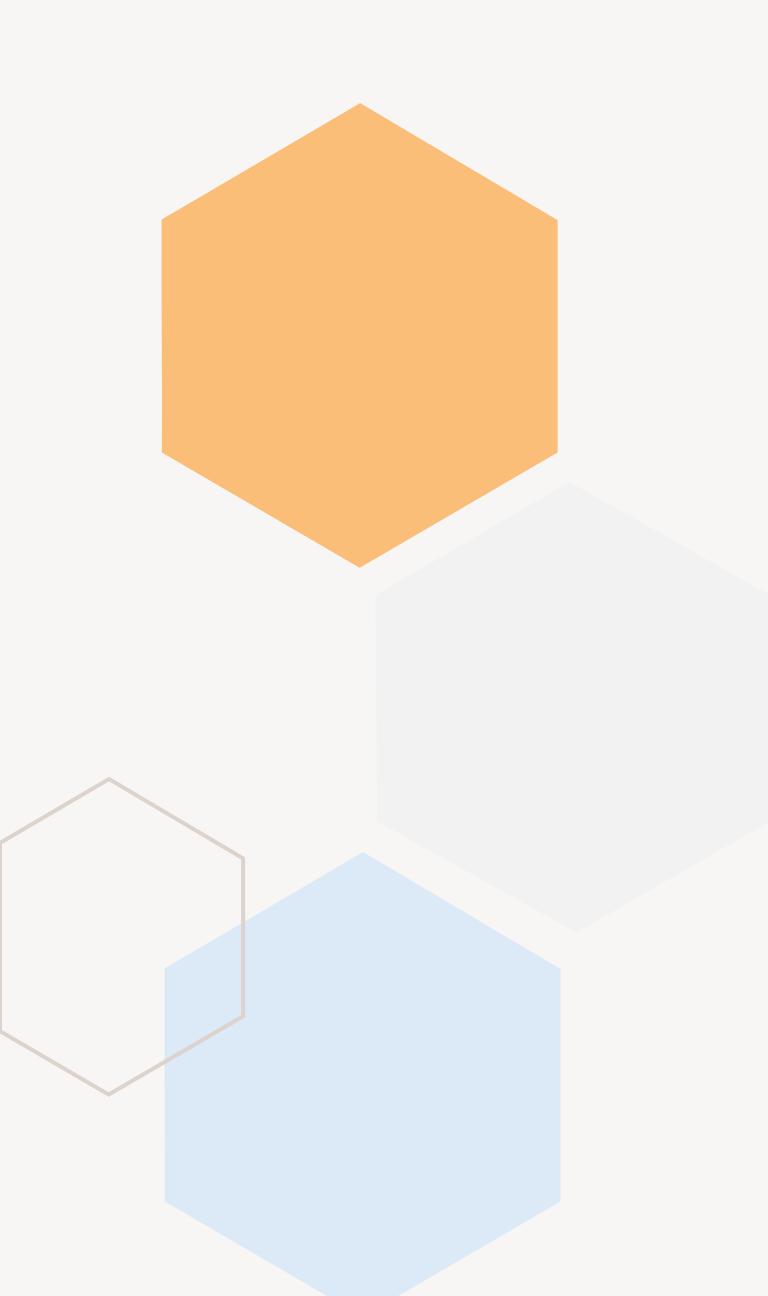
- Removed all pipe characters
- Removed tabs
- Re-generated clean PSV files

**Results:**  Success – Redshift COPY loaded all 7.7M records



# Final Architecture Diagram





# Key Takeaways

- ✓ Large datasets require careful memory management
- ✓ Weather data is highly incomplete
- ✓ ETL must handle outliers, missing values, and mixed formats
- ✓ Redshift COPY is strict — delimiter issues are common
- ✓ Iterative cleaning is essential for production-grade pipelines

# Phase 5: AWS S3 Upload & Region Migration



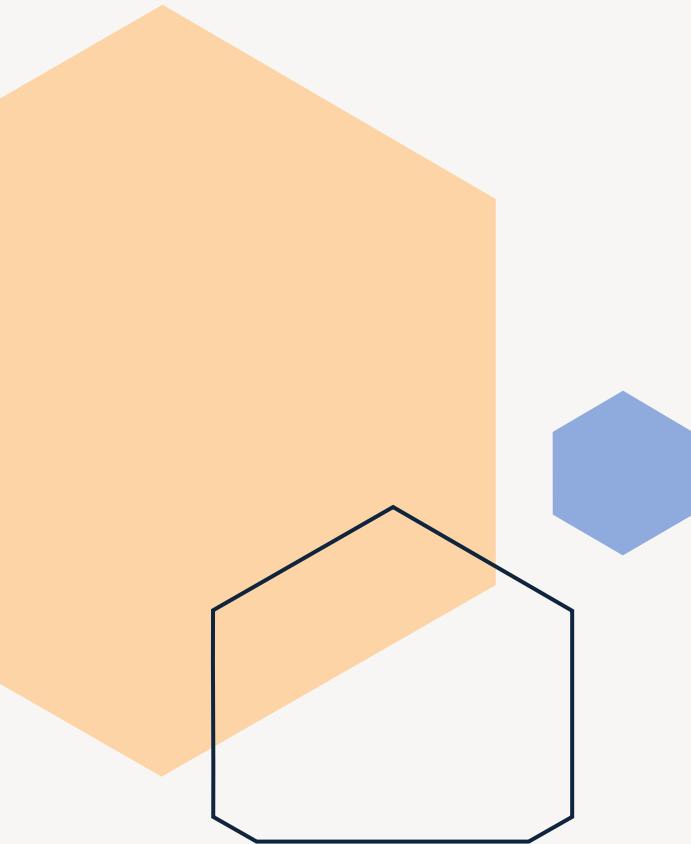
Goals: Move cleaned PSV files to AWS and prepare for Redshift COPY



# Phase 5: AWS S3 Upload & Region Migration

## Initial S3 Upload

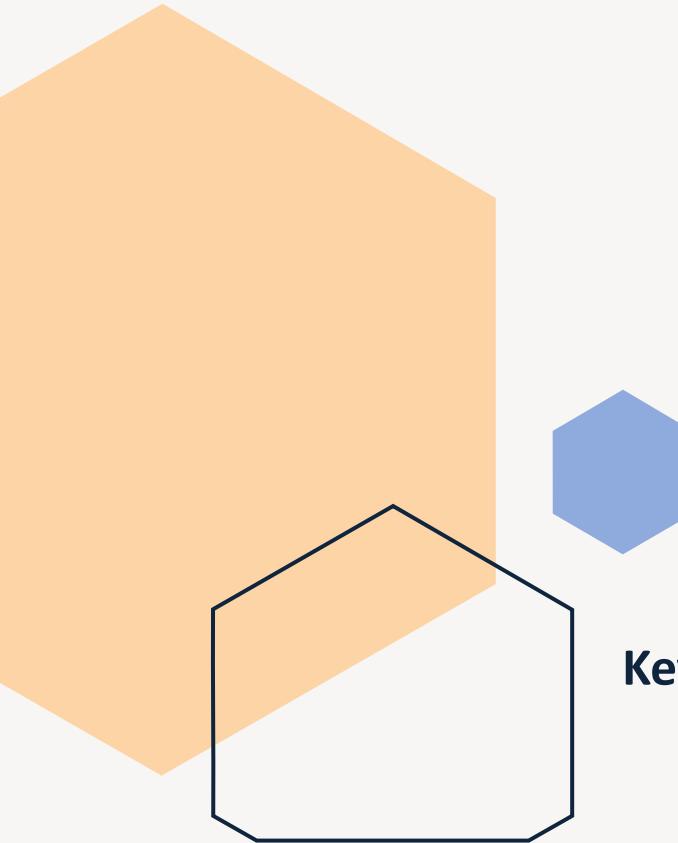
- Created bucket: us-accidents-data-project
- Region: **us-west-2 (Oregon)**
- Uploaded 5 CSV files
- Total upload size: ~2.1 GB
- Upload time: ~15 minutes



# Phase 5: AWS S3 Upload & Region Migration

## Region Migration

**Reason:** Lower latency + lower cost in us-east-1

- 
- ✓ Created new bucket: us-accidents-data-virginia
  - ✓ Region: **us-east-1**
  - ✓ Copied all files from us-west-2 → us-east-1
  - ✓ Verified file integrity
  - ✓ Deleted old bucket

**Key Insight:** Region choice matters for performance + cost

# Phase 5: AWS S3 Upload & Region Migration

## IAM Policy Update

### Action:

 Updated Redshift IAM role

 Added permissions:

- s3:GetObject
- s3>ListBucket

 Applied policy to Redshift cluster

### Outcome:

Redshift can now read from new S3 bucket

# Final Data Statistics

## Fact Table (fact\_accidents)

- Records: 7,728,234
- Format: PSV
- Size: ~1.8 GB

## Dimension Tables

| Table             | Records   | Size    |
|-------------------|-----------|---------|
| dim_location      | 2,847,562 | ~450 MB |
| dim_weather       | 1,456,789 | ~320 MB |
| dim_time          | 98,432    | ~8 MB   |
| dim_road_features | 4,567     | ~1 MB   |

## Total Volume

- ✓ 12,135,584 rows
- ✓ ~2.6 GB total
- ✓ S3 cost: ~\$0.50/month

# Data Quality Improvements

| Metric           | Raw     | Final | Improvement   |
|------------------|---------|-------|---|
| Duplicate IDs    | 0       | 0     | <span style="color: green;">✓</span> No cleaning Required |
| Null cities      | 253     | 251   | <span style="color: green;">✓</span> 99.% filled          |
| Weather outliers | 58,152  | 0     | <span style="color: green;">✓</span> 100% cleaned         |
| Format errors    | Unknown | 0     | <span style="color: green;">✓</span> Fully resolved       |
| Loading failures | N/A     | 0     | <span style="color: green;">✓</span> 100% success         |



# Key Lessons Learned (1/2)

## 1. CSV Format Challenges

- Redshift COPY is extremely strict
- Pipe delimiter + cleaning special chars solved it
- Always test COPY with small sample first

## 2. Missing Value Handling

- 13K missing cities broke location dimension
- Reverse geocoding + zipcode lookup fixed 99.6%
- Never assume data completeness

## 3. Outlier Detection

- Impossible values (150°F, 200mph winds)
- Domain knowledge required to set thresholds



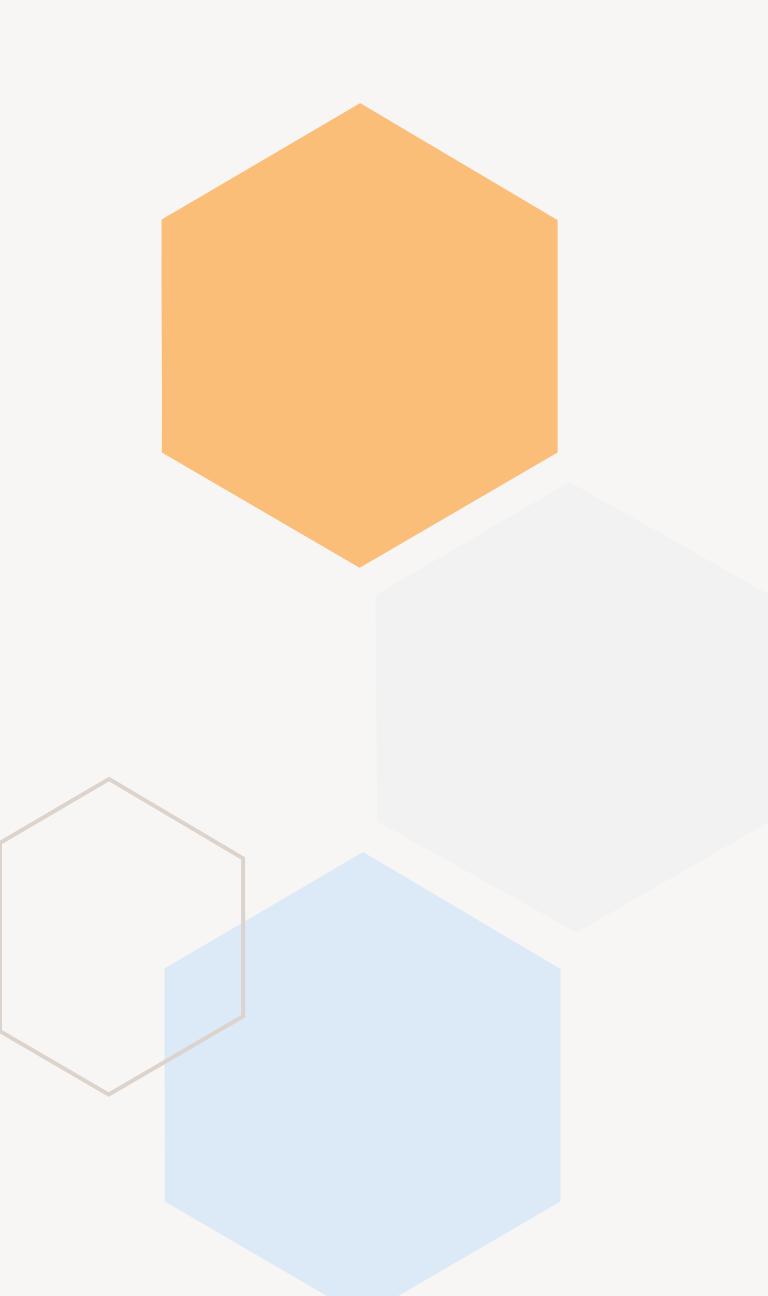
# Key Lessons Learned (2/2)

## 4. Iterative Development

- 9+ formatting iterations
- Upload → Test → Fix → Re-upload
- Real-world pipelines require resilience

## 5. Data Size Matters

- 3GB+ CSV requires careful memory management
- Chunking or powerful machine recommended
- Always check RAM before loading full dataset



# Tools & Libraries Used

## Python

- Pandas
- Numpy
- Geopy
- Uszipcode
- Boto3

## Environment

- Python 3.11
- Jupyter Notebook
- VS Code
- Windows 11

## AWS

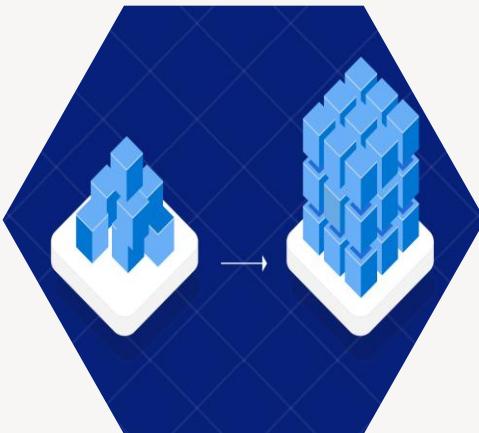
- S3
- IAM
- Redshift

# Data Architecture Approach



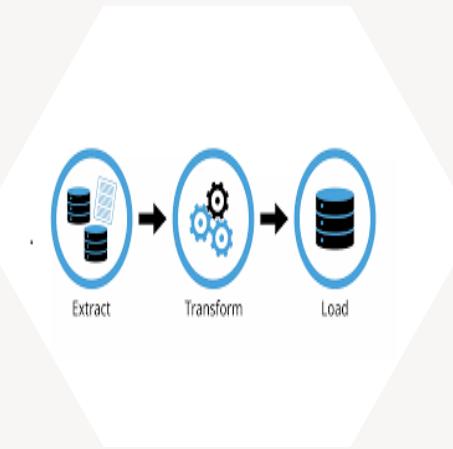
## Data Storage

Storing raw data in a structured database.



## Data Transformation

Clean, normalize, and preprocess data



## ETL Process

Extract, transform, and load the dataset for analysis and dashboarding



## Visualization

Build interactive dashboards

### Two major sections:

- Local Processing Pipeline (19 steps + ML)
- AWS Cloud Pipeline (S3 → Redshift → QuickSight )

# Conclusion

**What actually happened:**

- 4 iterations
- 13K cities imputed
- 58K outliers cleaned
- Star schema built
- Region migration completed
- Redshift COPY fully successful

**Data engineering is 80% cleaning, 20% analysis.**

**Key Insight:**

Real pipelines require iteration, validation, and resilience

# Thank you

